IPEDSBrain

Your IR Team's Research Analyst — On Demand

John Dimatos — February 2026

The Problem

Your IR office spends weeks producing a single peer comparison.

They download IPEDS files. Join tables in Excel. Manually compute metrics. Format reports. Build slides. Document methodology.

Then someone asks: **"Can you run that again with a different peer group?"**

And the work starts over.

The Problem Gets Worse Every Year

The tools are disappearing

NCES — the agency that maintains IPEDS — has been reduced to **3 staff members.** The free tools your IR office depends on may not be maintained.

The demands are growing

EDUCAUSE named "The Data-Empowered Institution" the **#1 IT issue for 2025.** Accreditors want more evidence. Boards want more data. The reporting burden only increases.

The methodology is fragile

When the Provost asks "how did you get that number?", the honest answer is often: *a chain of Excel operations that would take days to reconstruct*.

What If Your IR Team Could Do This

Ask a question in plain language:

> **"Give me a full profile of our institution with peer benchmarking against the Big Ten."**

And get back — in minutes, not weeks:

- A narrative report with every number cited to its federal data source
- Comparison tables showing exactly where you stand among peers
- A complete, auditable methodology document
- A script that any colleague can run to verify every number independently

Example: Institutional Profile

> `ipeds:profile University of Michigan–Ann Arbor`

What comes back:

Admissions & Selectivity

Admission rate: 18% (down from 23% in 2019). Among the 50 most selective public institutions nationally. Yield rate: 47%, ranking 3rd among Big Ten peers.
[Source: IPEDS DRVADM2023.DVADM01, UNITID 170976]

Student Outcomes

6-year graduation rate: 93%, up 2 points over 5 years. Ranks 1st among Big Ten peers and in the 96th percentile nationally among public universities.
[Source: IPEDS DRVGR2023.BAGR150, UNITID 170976]

Enrollment & Demographics

Example: Peer Benchmarking

```
ipeds:benchmark University of Michigan --peers big-ten
```

What comes back:

Metric	Michigan	Big Ten Median	Percentile	Rank
Admission rate	18%	57%	99th	1 of 14
6-yr graduation rate	93%	74%	99th	1 of 14
% Pell recipients	16%	24%	22nd	12 of 14
% URM enrollment	12%	15%	28th	10 of 14

Key finding: Michigan leads the Big Ten on selectivity and outcomes but lags on socioeconomic and racial diversity. Pell recipient share is 8 points below the conference median.

Example: Equity Audit

> `ipeds:equity University of Michigan`

What comes back:

Graduation Rate Gaps

Group	6-Year Grad Rate	Gap vs. Overall (93%)
White	95%	+2
Asian	96%	+3
Black	83%	**-10**
Hispanic	89%	-4
Pell recipients	87%	**-6**

Example: Trend Analysis

```
ipeds:trends admission_rate --portfolio big-ten --years 2019-2023
```

What comes back:

[LINE CHART: Big Ten Admission Rate Trends, 2019-2023]

Key findings:

- Michigan (-5 pts) and Northwestern (-7 pts) saw the steepest selectivity increases
- Rutgers (+1 pt) and Indiana (+2 pts) were essentially flat
- The Big Ten median admission rate dropped from 63% to 57%
- The spread between most and least selective widened from 48 to 54 points

> Trend analyses are automatically contextualized against the peer group and national benchmarks.

Everything Is Reproducible

Every analysis produces a self-contained bundle:

```
michigan-profile-2026-03-15.zip
├── README.md              How to reproduce this analysis
├── report.md              The full narrative report
├── analysis.py            Python script — reproduces every number
├── data/                  The exact data used (not the full IPEDS dump)
├── methodology.md         Every analytical choice documented
└── citations.json         Every claim traced to its source
```

Hand this zip to anyone. A colleague. An auditor. An accreditation reviewer.

They run one command: `pip install pandas && python analysis.py`

Every number in the report is verified. No AI required. No proprietary software. No API keys.

Why Reproducibility Matters

For accreditation

Reviewers can independently verify your evidence. Your methodology is documented, not buried in Excel.

For institutional trust

When the Provost asks "how did you get that number?", the answer is a 3-page methodology document and a script anyone can run — not "Sarah built a spreadsheet last year."

For continuity

Staff turn over. Institutional knowledge walks out the door. With IPEDSBrain, every analysis is a permanent, verifiable record of what was measured, how, and why.

What It's Not

It's not a dashboard

Dashboards show you data and wait for you to draw conclusions. IPEDSBrain answers questions — identifies patterns, flags gaps, contextualizes findings — and shows exactly how it got there.

It's not a data warehouse

No database to maintain. No ETL pipelines to monitor. No IT infrastructure to manage. It runs on a laptop.

It's not a black box

Every analysis generates a standalone script that reproduces every number without the AI. The methodology document explains every choice. The citation file traces every claim to a specific federal dataset.

Who Uses It

[DIAGRAM: Output Flow Through the Institution]

Produces	Consumes
IR Analysts — run analyses daily	**Provost** — institutional profiles, peer benchmarks
IR Director — define peers, set methodology	**Deans** — program-level analyses, demographic shifts
	Enrollment Management — yield data, competitor analysis
	Accreditation Coordinators — equity audits, outcomes evidence
	Board of Trustees — positioning slides, trend data

The Current Landscape

Your IR office probably uses some combination of:

What	Cost	Limitation
NCES free tools	Free	Manual, limited analysis, uncertain future
Tableau / Power BI	$5-60k/yr + staff time	Visualization only — no analysis, no narrative
Enterprise platforms (EAB, HelioCampus)	**$100k-$1M+/yr**	12-month implementation, vendor lock-in
Consulting firms (Huron, NCHEMS)	**$100k-$1M per project**	Expensive, slow, nothing left behind

IPEDSBrain: 8-12 week build. Runs on a laptop. No vendor lock-in. Every analysis is

IPEDSBrain

Ask a question. Get a cited, reproducible answer. In minutes.

Let's discuss how this fits your institution.